

# Evidence Before Effect: AETNAMEM’s Auditable Memory and Guarded-Action Control Plane for AI Agents

aetna000.com

Independent Researcher

aetna000@gmail.com <https://github.com/aetna000/aetnamem>

July 2026

Technical report; artifact version 1.0

## Abstract

Persistent memory and tool use turn language-model agents from text generators into stateful actors. The same transition creates two coupled failure channels: untrusted content can become durable “user memory,” and that memory can later be misused as authority for an external effect. We present AETNAMEM, a framework-neutral reference control plane that separates information from authority and binds memory transitions, action intent, approval, execution evidence, and outcome receipts into one independently verifiable history. Its memory path classifies origin, quarantines untrusted extractions, restricts retrieval to active records, preserves provenance, and records deletion and supersession. Its guarded-action path compiles proposed operations into a canonical, hash-bound `WORLDPATCH`; requires trusted, host-attested `authorized_by` evidence for enforced mutations; binds approval to the exact plan; revalidates adapter manifests and world-state preconditions; verifies postconditions; and records uncertainty instead of blindly retrying ambiguous effects. A per-subject SHA-256 chain and externally anchorable checkpoints support zero-dependency verification while erasable payload tables retain raw arguments and before-images.

We evaluate the open-source implementation at commit `Ocd082c`. A fresh deterministic run passes all 33 check-level assertions, all five scenarios, and all 81 severity-weighted points in `MEMORYSTACKBENCH`’s `seven_sins_v0_1` suite. This is a perfect conformance result, tied by eleven other targets in the pinned 19-target public snapshot; it is not a language-model score. The benchmark’s separate evidence matrix assigns AETNAMEM 16/18 auditability points, the highest value in that snapshot. A reproducible hostile-page demonstration then exercises five refusal paths, one authorized commit, two implementation-independent verifiers, and two tamper probes over a 12-event database. The results establish executable control-plane invariants for the tested paths, not universal prompt-injection immunity, formal verification, or forced mediation of tools outside the gate.

## 1 Introduction

Agent systems increasingly interleave model inference with retrieval, persistent memory, and tool calls. ReAct makes reasoning and acting explicit in one trajectory [24]; Toolformer demonstrates models that select and invoke APIs [19]; and memory architectures such as Generative Agents, Mem-

oryBank, and MemGPT retain information across interactions [16, 25, 15]. These capabilities are useful precisely because they move information and alter state. They are dangerous for the same reason.

The usual agent diagram ends at “tool result.” That endpoint is too early for a dependable system. A tool call may have lacked authority, executed a stale plan, raced with another writer, returned success without changing the world, changed the world before timing out, or left a result that cannot later be attributed to the approved request. Persistent memory adds a second delayed channel: a hostile webpage can influence a model in one session, become stored as a user fact, and authorize or bias an effect in a later session. Indirect prompt injection is therefore not only a prompt-level problem [8, 5]. It is an information-flow, authority, transaction, and evidence problem.

Classical systems security offers durable principles: fail-safe defaults, complete mediation, separation of privilege, and least privilege [18]; capability-based authority [6, 13]; compensating transactions for effects that cannot be atomically rolled back [7]; provenance that explains where data came from [3]; and tamper-evident logs anchored outside the system they describe [20, 9]. The central thesis of this work is that an agent control plane should apply those ideas around a fallible model rather than expecting the model to emulate them reliably in natural language.

AETNAMEM is a compact reference implementation of that thesis. It is not another agent loop and does not depend on a particular model, planner, or host. It supplies policy-bearing memory, guarded external actions, and a shared evidence ledger that can sit beneath OpenClaw, Hermes-like hosts, LangGraph, Letta, or any process able to call Python, a CLI, or MCP. The language model may propose; the control plane decides whether the proposal can become durable state or an external effect.

### 1.1 Design question

We ask:

*Can a small, framework-neutral systems layer preserve useful agent memory and tool access while ensuring that untrusted information does not silently become authority, approved intent cannot drift before execution, ambiguous effects remain explicit, and historical claims are independently checkable?*

This question is narrower than “solve agent safety.” It is also more operational. The desired output is not merely a refusal

message. It is a set of durable artifacts: source classification, policy decision, exact plan, approval scope, before-state evidence, execution attempt, observed after-state, receipt, and a chain position that another implementation can verify.

## 1.2 Contributions

This report makes five contributions.

1. **Information-authority separation.** Evidence references carry both a causal relation (`informed_by` or `authorized_by`) and a trust tier. Untrusted content may shape a proposal but cannot authorize an enforced mutation.
2. **Trust-aware persistent memory.** Webpage, tool-output, and derived facts enter quarantine with provenance; only active records are recallable; and quarantined values cannot supersede trusted facts until explicit promotion.
3. **Exact-plan guarded actions.** Canonical `WORLDPATCH` documents bind arguments, previews, preconditions, adapter manifests, dependencies, evidence, policy, and identity context into a plan hash. Approval and terminal receipts bind to that exact plan.
4. **Dual-plane, independently verifiable evidence.** Raw erasable payloads are separated from digest-only append history. Per-subject hash chains, checkpoints, deletion receipts, action receipts, and standard-library verifiers make database claims testable without importing `AETNAMEM` code.
5. **Reproducible evaluation artifacts.** We report a pinned benchmark snapshot, a fresh current-version rerun, the complete 88-test repository suite, and a deterministic hostile-page/action demonstration whose SQLite database and checkpoints are distributed with the code.

## 1.3 Headline result, stated precisely

The `AETNAMEM` adapter achieves 33/33 checks and 5/5 scenarios on `MEMORYSTACKBENCH` [12]. In colloquial terms it “aced” the suite. Scientifically, three qualifiers matter. First, the target is a deterministic engine with no language model, so this result does not measure model intelligence. Second, eleven other targets also reach 33/33 in the pinned snapshot, so this is perfect conformance rather than a unique win. Third, the same GitHub account maintains `AETNAMEM` and `MEMORYSTACKBENCH`; we therefore describe the run as reproducible self-evaluation, not independent third-party validation. These qualifiers do not weaken the executable result; they define what it means.

The remainder of the paper defines the threat model, presents the architecture and protocols, reports evaluation results, and separates implemented guarantees from deployment assumptions and roadmap work.

# 2 Scope, Terminology, and Threat Model

## 2.1 System boundary

We use *agent* for the model-driven process that interprets a task and proposes memory or tool operations. We use *host* for the process that knows channel identity, origin, session

context, and deployment policy. A *reviewer* holds approval authority outside the agent-facing process. An *adapter* translates a typed operation into a provider-specific effect and observation. The *world* is the mutable external resource, such as a filesystem, mailbox, browser, cloud API, or database.

`AETNAMEM` is a control plane, not a model. It can run below a deterministic script, an LLM agent, or a human-operated CLI. In the current implementation, memory tools are available through MCP and a native OpenClaw plugin, while guarded actions are exposed through Python and CLI interfaces. A universal action-interception MCP gateway remains roadmap work.

## 2.2 Assets and adversaries

The protected assets are: (i) trusted active memory; (ii) user data named in memory and action payloads; (iii) external resources reachable through adapters; (iv) reviewer approval authority; and (v) the integrity and continuity of the evidence history.

We consider an adversary who can:

- place arbitrary instructions in webpages, documents, tool output, or assistant-generated text that an agent processes;
- induce the agent to propose an operation, cite malicious content as evidence, omit authority, or modify a proposal after review;
- modify mutable SQLite rows after staging or after execution;
- edit, insert, reorder, or delete local audit rows and replace local artifacts;
- trigger exceptions or process interruption around an external-effect boundary; and
- race external resource state between proposal and commit.

The base analysis assumes the adversary cannot break SHA-256 collision resistance or HMAC unforgeability, obtain the reviewer key, alter a genuinely external checkpoint anchor, or compromise the trusted host’s source/identity attestation. Compromise of those roots of trust is discussed as residual risk.

## 2.3 Security goals

## 2.4 Non-goals

The current release does not claim to: detect every prompt injection; authenticate a caller that lies about `source_type`; force an agent through the gate when it also has direct write tools; provide a hardened operating-system sandbox; prove adapter correctness; offer non-repudiation from a shared HMAC key; make SQLite bytes disappear from free pages, WAL files, backups, or replicas; provide Byzantine consensus; or establish legal/regulatory compliance. It also does not claim statistical generalization from five deterministic benchmark scenarios. These boundaries are part of the design contract, not footnotes to it.

## 2.5 Trusted computing base

The trusted computing base (TCB) consists of the source classifier or host attestation, deterministic policy functions, canonical serializer and hash primitives, approval-key custody, adapter implementation, SQLite/OS behavior, verifier

Table 1: Security goals and their concrete enforcement points.

| Goal                    | Required behavior                                              | Enforcement point                                             |
|-------------------------|----------------------------------------------------------------|---------------------------------------------------------------|
| G1 Origin integrity     | Preserve source class and provenance for every candidate fact  | Host/classifier, record schema, audit event                   |
| G2 Quarantine           | Prevent untrusted facts from recall or trusted supersession    | Admission policy, active-only retrieval                       |
| G3 Authority            | Prevent informational evidence from authorizing mutation       | Evidence relation, trust tier, attestation check              |
| G4 Exact intent         | Execute only the reviewed operational plan                     | Canonical plan hash, approval signature, commit recomputation |
| G5 Fresh world state    | Abort or compensate when assumptions change                    | Before-image preconditions and revalidation                   |
| G6 Outcome evidence     | Distinguish requested, attempted, observed, and verified state | Adapter receipt, postcondition verification, terminal receipt |
| G7 Historical integrity | Detect local history manipulation relative to a trusted anchor | Hash chain, checkpoint, standalone verifier                   |
| G8 Erasure boundary     | Purge live content without erasing structural accountability   | Mutable payload plane, digest-only audit plane                |

implementation, and external checkpoint sink. The agent and untrusted content are explicitly outside the TCB. A model may be fully prompt-injected and still encounter deterministic denials, provided complete mediation holds at deployment.

This decomposition echoes the instruction/data separation pursued by instruction-hierarchy training [22] and CaMeL’s control/data-flow separation [4], but moves the enforcement target into persistent memory and transaction evidence. Unlike CaMeL, AETNAMEM does not presently claim a formal information-flow proof or universal mediation.

### 3 Architecture

Figure 1 shows two policy paths sharing one evidence substrate. The upper path governs what becomes durable and recallable. The lower path governs what may alter the world. Both paths emit causal metadata and digests into a per-subject event chain, while raw content and before-images remain in mutable tables with explicit retention semantics.

#### 3.1 Control-plane decomposition

The architecture is organized into seven logical stages.

**1. Origin-aware ingestion.** The host supplies or derives source class, subject, session, and turn. The extractor proposes facts but does not choose their authority. In the reference extractor, user statements and embedded webpage/tool payloads are deterministically distinguished.

**2. Memory policy.** Candidate facts receive a trust tier and initial lifecycle state. Trusted user facts become active. Untrusted or derived facts become quarantined and retain their origin, episode, confidence, slot, session, and turn.

**3. Active-only context construction.** Recall ranks only active records and writes candidate/score evidence. Persona and context blocks are live-derived, bounded, and annotated with source record identifiers, reducing stale cached summaries and feedback loops.

**4. Intent compilation.** An action proposal is normalized by an adapter into typed arguments, preview, effect class, preconditions, sensitive-field declarations, and a manifest fingerprint. Dependencies and evidence are ordered canonically. The resulting WORLDPATCH is content-addressed.

**5. Authority and approval.** Policy checks evidence before a plan can enter an executable state. A separate reviewer signs the exact plan hash with an expiry and nonce. The key is absent from the agent-facing process.

**6. Scoped execution and verification.** Commit recomputes the persisted plan, checks approval, verifies adapter identity, revalidates all preconditions, then crosses the external boundary. The adapter observes after-state and verifies the postcondition. Ambiguity is represented as uncertainty, not success.

**7. Evidence admission.** Every transition emits a structural event. A receipt binds terminal state and operation digests to the plan and event chain. Standalone verifiers recompute these relationships. Checkpoints export chain heads to a separate trust domain.

#### 3.2 Three planes, not one log

A single append-only log cannot simultaneously maximize audit integrity, operational utility, privacy, and erasure. AETNAMEM therefore separates:

**Live state plane.** Episodes, records, retrieval details, action transactions, and status needed for normal operation.

**Sensitive payload plane.** Raw action arguments, previews, before-images, provider results, compensation material, and optionally retained query text. This plane is explicitly purgeable.

**Evidence plane.** Event types, identities as claimed/attested, record and operation IDs, hashes, transitions, and receipt bindings. Engine-generated events avoid raw fact values and secrets.

This is not full cryptographic privacy. Digests can leak equality and are vulnerable to dictionary attacks when the underlying value has low entropy. The split instead provides a tractable retention boundary: structural accountability can survive after operational payloads are purged.

#### 3.3 Framework neutrality

Framework neutrality means the policy semantics do not depend on the agent loop. The same core can be reached through Python, CLI, stdio MCP, or a host-native lifecycle adapter. A framework integration should remain thin:

- before model context assembly, request a bounded persona/recall block;
- after a turn, capture user facts and log assistant/tool content as classified data or digests;
- for a mutation, submit a typed proposal and causal evidence rather than a raw shell handle;
- keep reviewer authority in another process; and
- retain audit identifiers in the host’s own trace.

This pattern allows a host to replace its model or planner without silently replacing memory and action semantics.



active records with the same slot. A quarantined record never triggers supersession. This asymmetry matters: a hostile page can target exactly the same slot as a trusted user fact and still cannot replace it.

Trusted candidates deduplicate only against active records. Otherwise, a previously quarantined copy of the same text could swallow a later genuine user statement and keep it unavailable. Quarantined candidates may deduplicate against active or quarantined records to limit redundant evidence while preserving the active version.

**Property M2 (Trusted-slot preservation).** For an active record  $m$  and an incoming untrusted candidate  $c$  with the same fact slot, admission of  $c$  leaves  $m$  active and places  $c$  in quarantine.

The property is a direct consequence of the supersession condition: old records are selected only when the new candidate status is active. The flagship experiment targets this exact edge case.

## 4.5 Promotion and derived memory

Promotion is explicit, subject-scoped, and audited. It changes status to active, raises trust to `user_confirmed`, and only then applies same-slot supersession. Deployments that require human-only promotion must not expose the promotion tool to the agent without another approval boundary.

Higher layers are derived rather than independently authoritative:

- **L0 episodes:** raw ingested turns, subject to purge;
- **L1 records:** extracted facts with lifecycle and provenance;
- **L2 scenes:** deterministic session-grouped views;
- **L3 persona:** live-derived bounded view of active records; and
- **proposals:** external synthesis results quarantined pending admission.

Because persona and scenes are generated from active L1 state, they do not persist an independent stale copy. Generated context carries record IDs and emits its own audit event.

## 4.6 Deletion semantics

Forget operations match both active and quarantined records. The engine tombstones records, clears content and fact slots, removes FTS entries, purges source episodes that no longer support live records, and returns a deletion receipt. The receipt binds selector digest, purged IDs, event ID, and event hash. This is a logical purge of the live database, not physical sanitization of storage media or downstream copies.

Deletion applies to quarantine because otherwise an injection payload could remain in a hidden review queue after the user requested erasure. It also produces evidence because an unverifiable “deleted” response is not enough for a durable privacy workflow.

# 5 Guarded Actions

The action engine treats tool mutation as a durable causal transaction. The model may propose arbitrary text, but only a typed adapter can prepare an executable operation, and only an exact approved plan can cross the external-effect boundary.

## 5.1 Evidence is not authority

Each evidence reference is a tuple

$$e = (\text{kind}, \text{ref}, \text{digest}, \text{relation}, \text{trust}, \text{attested}), \quad (4)$$

where relation is either `informed_by` or `authorized_by`. The distinction is structural, not prompt wording. A webpage may legitimately inform an operation (for example, supply a shipping quote) while lacking authority to change the user’s delivery address.

For a prepared operation  $o$ , evidence set  $E$ , and mode  $\mu$ , the core policy is:

$$\begin{aligned} d(o) = \text{unknown\_blocked} &\implies \neg \text{allow}(o, E, \mu), \\ \text{mut}(o) \wedge A(E) = \emptyset &\implies \neg \text{allow}(o, E, \text{enforce}). \end{aligned} \quad (5)$$

Here  $\text{mut}(o)$  denotes a mutating effect, and

$$A(E) = \{e \in E \mid e.\text{relation} = \text{authorized\_by}, e.\text{trust} \in T, e.\text{attested}\}, \quad (6)$$

and  $T = \{\text{trusted\_user}, \text{user\_confirmed}\}$ . Any purported authority outside  $T$  is rejected even when another valid authority exists; this avoids silently accepting contaminated authority sets.

**Property A1 (Non-escalation of information).** An evidence reference marked `informed_by` can never satisfy the enforced-mutation authority requirement. An `authorized_by` reference with untrusted tier or absent host attestation is rejected.

## 5.2 Canonical WorldPatch

An operation proposal names an adapter, typed operation, arguments, causal evidence, and dependencies. Adapter preparation returns normalized arguments, preview, preconditions, effect class, and sensitive fields. The engine computes digests for arguments, preview, preconditions, and adapter manifest. It topologically orders operations and rejects unknown dependencies or cycles.

Let  $J(x) = \text{Canon}(x)$  denote UTF-8 JSON with sorted keys and no insignificant whitespace, and  $H(x) = \text{SHA256}(J(x))$ . The plan hash is

$$p = H(\text{format}, tx, u, a, \mu, v, H(\text{policy}), [o_1, \dots, o_n]), \quad (7)$$

where every  $o_i$  contains operation identity, effect class, argument/preview/precondition/manifest digests, dependency IDs, and sorted evidence. Transaction, subject, actor, mode, plan version, and policy are therefore in approval scope.

## 5.3 Approval

In `enforce` mode, a valid proposal enters `awaiting_approval`. A separate process issues

$$\sigma = \text{HMAC}_k(J(tx, p, \text{approver}, t_i, t_e, n, \text{format})), \quad (8)$$

using SHA-256 HMAC [11]. The document contains issue/expiry time and a 128-bit random nonce. The database

enforces nonce uniqueness. The shared key demonstrates possession of  $k$ , not the real-world identity of the free-text approver label. KMS-backed asymmetric signers can implement the same issue/verify surface in future deployments.

**Property A2 (Exact-plan approval).** Under HMAC unforgeability and hash collision resistance, an approval for plan hash  $p$  cannot authorize a different plan  $p'$  unless the verifier is bypassed, the key is compromised, or a relevant hash collision is found.

## 5.4 Commit protocol

Before the first effect, commit performs the following checks:

1. transaction mode is `enforce` and state is `approved`;
2. the persisted plan recomputes to the recorded hash;
3. a matching proposed event exists on the subject chain;
4. approval digest, signature, scope, and expiry are valid;
5. current adapter manifests match approved manifest digests; and
6. every operation's preconditions still hold.

The engine then durably enters `committing`. Immediately before each operation it revalidates preconditions again, records `effect_executing`, and only then invokes the adapter with a deterministic idempotency key. The adapter returns a provider request ID, result, and observed after-state. A separate adapter method re-reads authoritative state and verifies the postcondition.

## 5.5 Effect classes

Adapters classify operations as:

**read only** no mutation;

**transactional reversible** provider supplies atomic rollback semantics;

**verified compensatable** before-state can be restored and checked, but external observers may have seen the effect;

**irreversible staged** execution requires explicit staging and cannot be generically undone;

**observe only** no execution; or

**unknown blocked** fail closed.

The filesystem reference adapter is `verified_compensatable`, not `transactional`. It confines relative paths to a configured root, rejects symlinks and non-regular files, captures content hash/mode/before-image, atomically replaces file contents, and verifies observed existence/hash/mode. Filesystem hooks and observers can still see intermediate effects; compensation is not time reversal.

## 5.6 Failure and uncertainty

No SQLite transaction remains open while a provider call runs. If dispatch raises, the system cannot generally know whether the provider applied the effect before failing. The operation and transaction become `uncertain`; already verified earlier operations are compensated where possible. A process that dies in `committing` is fenced by recovery into `recovery_required`. Generic recovery never blindly retries. Provider-specific idempotency lookup or authoritative read must resolve the outcome.

**Property A3 (No success by assertion).** A transaction reaches `committed` only after every operation's postcondition verifier returns true. Provider return alone is insufficient.

**Property A4 (No generic blind retry).** An interrupted in-flight effect is represented as uncertain/recovery-required until provider-specific reconciliation. This avoids turning an ambiguous first execution into an unintended duplicate.

These are state-machine invariants of the reference implementation, not guarantees that an adapter's observation is semantically complete.

# 6 Evidence Ledger and Independent Verification

The evidence ledger makes claims falsifiable after execution. It does not prevent all compromise; it provides a cryptographic consistency relation between events, plans, receipts, and externally pinned history.

## 6.1 Per-subject event chain

Each event stores sequence, event ID, subject, type, UTC time, actor, optional session/turn/record, canonical payload, previous hash, and event hash. For subject  $u$ ,

$$h_i = H(\text{id}_i, u, \text{type}_i, t_i, a_i, s_i, r_i, \text{payload}_i, h_{i-1}), \quad (9)$$

with  $h_0 = \perp$ . Global sequence numbers order database insertion, while chain links are per subject. The verifier recomputes every event preimage and checks continuity.

Engine-generated memory events contain message/query-selector digests and structural metadata rather than raw values. Action events similarly carry plan, result, observation, and error digests while raw payloads live elsewhere. Low-level custom action logging can still accept arbitrary caller payloads, so deployment code must preserve this convention.

## 6.2 Checkpoints

A checkpoint stores, for every subject, the current sequence, head hash, and event count, then hashes the checkpoint document itself. Verification applies a containment rule: the pinned sequence must exist and have the pinned hash. A checkpoint must be copied to a trust domain the database owner cannot rewrite, such as WORM object storage, a transparency service, or an RFC 3161 timestamp authority [1]. A checkpoint left beside the database mainly detects accidents.

Hash chaining detects editing, insertion, reordering, and middle deletion when the verifier has a trusted expected head. It cannot by itself detect deleting a suffix or replacing the entire database and recomputing every hash. The external head closes that gap relative to its anchoring time.

## 6.3 Receipts

A deletion receipt binds subject, time, selector digest, purged record/episode IDs, audit event ID/hash, and receipt digest. An action receipt binds transaction, subject, exact plan hash, terminal state, operation-result/compensation digests, audit event ID/hash, and receipt digest. Verification requires:

1. the receipt digest recomputes;
2. its event exists on a valid subject chain;
3. the event hash equals the receipt's bound hash;




Figure 2: Guarded-action lifecycle. Database transactions protect ledger transitions, not external effects. The engine records execution intent before crossing the adapter boundary and represents ambiguous outcomes explicitly as *uncertain* or *recovery required*.

4. event payload and receipt agree on transaction, plan, terminal state, and operation digest; and
5. plan and recorded operational tables recompute consistently.

## 6.4 Independent implementations

The repository ships two Python-standard-library programs. `verify_audit.py` implements the event-chain, checkpoint, and deletion-receipt rules without importing `AETNAMEM`. `verify_actions.py` additionally reconstructs `WORLD-PATCH`, validates approval scope and optional HMAC, and checks action receipt bindings. This is implementation independence: a defect in the engine’s self-check need not be repeated by calling the same library. It is not organizational independence; both programs are maintained in the same project.

**Property E1 (Mutation detection relative to an anchor).** If an adversary modifies an event at or before a trusted checkpoint’s pinned head without finding a hash collision, then either the chain rule or checkpoint containment check fails.

*Argument.* Modifying event  $e_j$  changes  $h_j$  by Equation (9). Retaining the old suffix breaks the next `prev_hash`; recomputing the suffix changes the pinned head. Deleting the suffix fails containment. Replacing the external checkpoint lies outside the assumption.

## 6.5 Evidence is not truth

The ledger proves consistency among recorded assertions. It does not make a false adapter observation true, authenticate a forged host source label, or prove that no unlogged side channel existed. In particular:

- *integrity* means recorded bytes have a consistent hash history;
- *authenticity* depends on keys, host identity, and adapter/provider evidence;
- *completeness* depends on mediation and event coverage; and
- *semantic correctness* depends on policy and verifiers.

This distinction avoids treating “blockchain-like” logging as a substitute for access control or trustworthy observation.

## 7 Reference Implementation

The evaluated artifact is `AETNAMEM v0.2.1`, an AGPL-3.0 Python package requiring Python 3.10 or newer [2]. Runtime dependencies are limited to the Python standard library and SQLite. Development adds `pytest`. The checkout contains approximately 9,380 lines across 42 tracked Python source/test/tool files at the evaluated commit.

### 7.1 Storage and interfaces

SQLite stores episodes, records, FTS5 indexes, retrieval events, audit events, action transactions, operations, dependencies, evidence, approvals, attempts, payloads, and receipts. Python API and CLI expose memory operations, inspection, checkpointing, and guarded actions. The stdio MCP server exposes memory verbs through newline-delimited JSON-RPC. A native TypeScript OpenClaw plugin performs pre-prompt recall and post-turn capture using the same engine.

The action subsystem currently exposes a Python protocol and CLI. Its first adapter supports root-confined UTF-8 `write_text` and `delete_file`. Additional providers and a mandatory mediation gateway are not yet implemented.

### 7.2 Atomicity boundary

The implementation uses SQLite transactions to atomically couple live-state changes and corresponding audit events. For example, a memory record and its `record_created` event succeed or roll back together; deletion, payload purge, and their evidence follow the same rule. External effects are deliberately outside database transactions. The coordinator commits intent before invocation, then records observation and terminal state after invocation.

This boundary avoids holding a database lock across a slow API, but introduces an unavoidable uncertainty window. The state machine makes that window explicit and recoverable rather than pretending distributed effects share SQLite atom-icity.

### 7.3 Canonicalization and identifiers

Hashed structures use sorted-key compact JSON, UTF-8, and lowercase SHA-256 hex. Record, event, transaction,




Figure 3: Evidence binding in the recorded flagship run (intermediate events are visually elided, not omitted from verification). Hash chaining detects prefix mutation; an externally held checkpoint is required to detect whole-history replacement or suffix truncation. The construction is tamper-evident, not tamper-proof.

operation, evidence, approval, attempt, and receipt identifiers use typed prefixes and random UUID/nonce material. Idempotency keys derive from transaction/operation identity, adapter/operation, and argument digest.

Canonicalization is versioned through document format strings. The action plan, approval, receipt, and checkpoint formats each carry a distinct v1 identifier. Changing a hashed schema requires a new format version rather than silent reinterpretation.

## 7.4 Policy compactness

The critical memory policy is intentionally small: source classification, source-to-trust mapping, initial status, slot-aware supersession, deduplication, and deletion intent. The action policy is similarly narrow: unknown effects fail closed; enforced mutations require at least one trusted, attested authority; and untrusted authority is rejected. These functions contain no benchmark-specific vocabulary and are directly tested with unrelated examples.

Compact policy does not imply a small product surface. It means the decision kernel can be read and tested separately from extraction, ranking, host integration, and adapters.

## 7.5 Artifact layout

The repository distributes:

- machine-readable audit and action specifications;
- standalone verifier programs;
- 88 unit/integration/documentation tests;
- benchmark adapter and target manifest;
- a deterministic flagship driver with hostile HTML fixture;
- an actual SQLite database, checkpoint JSONL, and transcript from the recorded run; and
- this paper’s pinned CSV data, plotting script, diagrams, evidence validator, and LaTeX source.

The paper build does not scrape the live leaderboard. This prevents an updated website from silently changing a submitted figure. Data provenance records source commit and SHA-256 digests; refresh is a deliberate operation.

# 8 Evaluation Methodology

The evaluation answers four research questions.

**RQ1:** Does the memory engine satisfy the tested write, retrieval, deletion, untrusted-source, and temporal-update contract?

Table 2: MemoryStackBench suite composition.

| Scenario                             | Category                    | Checks | Severity           |
|--------------------------------------|-----------------------------|--------|--------------------|
| Durable private-itinerary preference | Retrieval correctness       | 6      | Medium             |
| Webpage memory poisoning             | Untrusted-source resistance | 10     | High               |
| SFO → OAK correction                 | Temporal update handling    | 7      | Medium             |
| Backup-email deletion                | Deletion behavior           | 5      | High               |
| Avoiding beef ⇒ vegetarian           | Write correctness           | 5      | Medium             |
| Total                                |                             | 33     | 15 high, 18 medium |

**RQ2:** Does the action path refuse missing/untrusted authority and plan drift while permitting an exactly approved effect?

**RQ3:** Can independent code paths verify the resulting database and detect concrete plan/history tampering?

**RQ4:** What evidence surfaces does the public benchmark observe relative to other targets?

## 8.1 Pinned artifacts and environment

The evaluated engine is v0.2.1 at Git commit 0cd082c9cac14f35a66ff946395a31847322005d. The benchmark code and public comparison snapshot are pinned at 10b9407ce54c92bcb8aee24099505427aabeebcd. We executed the experiment on macOS 26.5.1 (build 25F80), Apple M2, 16 GiB memory, Python 3.12.2, and SQLite 3.50.6 on 13 July 2026.

All benchmark and flagship tasks are deterministic. No LLM, API key, network request, embedding service, or stochastic sampler participates in the AETNAMEM run. Consequently we report exact conformance counts, not means or confidence intervals. A repeated run tests reproducibility of the implementation, not sampling variance.

## 8.2 MemoryStackBench

MEMORYSTACKBENCH executes a common adapter interface over multi-session *plant* → *interfere* → *probe* scenarios. The *seven\_sins\_v0\_1* suite contains five scenario-level tests and 33 assertions:

Checks inspect both user-visible responses and normalized memory records. They test required provenance fields, absence of poisoned/deleted/stale/overgeneralized content, and presence of the intended active fact. Severity weighting assigns 81 total points in the pinned suite.

We performed two analyses:

1. a fresh run against the current AETNAMEM checkout via PYTHONPATH; and
2. a descriptive comparison using the public 19-target leaderboard and auditability JSON files, whose hashes are recorded in the artifact.

The public targets are heterogeneous. Some are native semantic memory systems; others are store-level or session-level harnesses with benchmark adapter policy. Scores therefore establish conformance of each configured target, not a universal ranking of every product capability. The benchmark itself states this distinction. Timing logs are informational and are not used for performance claims.

### 8.3 Auditability matrix

The benchmark separately scores six evidence dimensions from zero to three: inspectability, provenance, retrieval transparency, deletion evidence, mutation lineage, and tamper evidence. It also records whether evidence is declared native, adapter-injected, or undeclared. This is a descriptive matrix, not part of the 33 safety checks.

For AETNAMEM, all six origins are declared native. Other rows in the pinned public snapshot have undeclared origin metadata. Cross-target point differences should therefore be interpreted as surfaces exposed to the harness, not independently audited security strength.

### 8.4 Flagship hostile-page experiment

The deterministic driver executes three acts, illustrated in Figure 4.

**Act 1: delayed memory poisoning.** A trusted user fact establishes `report.md`. A benign-looking HTML fixture contains a hidden instruction that targets the same fact slot with `files.attacker.example/steal.md`. A naive HTML text extraction includes the payload. The engine receives it as webpage content, extracts the candidate, and must quarantine it. Recall must return only the trusted path.

**Act 2: authority and exact-plan enforcement.** The driver attempts five prohibited paths: untrusted memory as action authority; mutation with no authority; commit without reviewer key; commit of an unapproved transaction; and commit after mutation of a persisted approved plan on a copy. Every prohibited step is wrapped in a shell assertion that fails the demo if the command exits successfully. A genuine host-attested task is staged, approved with an expiring HMAC in a separate process context, committed through the filesystem adapter, and verified against observed state.

**Act 3: independent evidence checks.** The engine exports a checkpoint and verifies itself. Then the two standalone programs verify the database. Finally, a copy’s quarantine event is rewritten from quarantined to active; the audit verifier must exit nonzero and identify sequence 4.

The recorded artifact includes the real `memories.db`, checkpoint JSONL, and transcript. Our reproduction reran the driver in a temporary directory and separately re-executed both standalone verifiers against the recorded copies.

## 8.5 Repository tests

We ran the complete test suite, not only benchmark tests. It covers memory transactions, policy, auditability, retrieval, extraction, MCP, CLI, host adapters, layers, journal import, actions, and documentation. Action tests include wrong-plan approval, nonce replay, world-state race, adapter drift, payload purge, operational-table tamper, uncertain execution, recovery fencing, and failed compensation.

## 8.6 Validity controls

To reduce self-confirming evaluation:

- benchmark vocabulary does not appear in core policy functions;
- direct unit tests use non-benchmark facts and resources;
- refusal steps assert nonzero exit rather than matching a favorable transcript alone;
- standalone verifiers do not import the package under test;
- tamper probes operate on copies so the distributed original remains verifiable; and
- source commits, environment, data hashes, and exact commands are disclosed.

Remaining construct and external-validity threats are addressed in Section 12.

## 9 Results

### 9.1 RQ1: memory conformance

The fresh current-version run passes 33/33 check-level assertions, 5/5 scenarios, and 81/81 severity-weighted points. Category results are complete: write correctness 5/5, retrieval correctness 6/6, deletion behavior 5/5, untrusted-source resistance 10/10, and temporal-update handling 7/7. There are no failure records.

Figure 5 places that result in the public comparison. Twelve of 19 targets reach 33/33; two reach 32/33; two reach 31/33; and the remaining three score 30, 26, and 24. Thus the correct claim is *perfect conformance in a twelve-way tie*. It would be inaccurate to describe the result as a unique state of the art.

The category matrix in Figure 6 shows where non-perfect targets lose checks. Untrusted-source resistance accounts for at least one failure in every target below 33/33. The weakest public row also loses retrieval, temporal-update, and write checks. This descriptive pattern supports the design focus on provenance and trust gating, but the small suite is insufficient for causal attribution.

### 9.2 RQ4: auditability evidence

AETNAMEM receives 16/18 points: 3/3 for inspectability, provenance, retrieval transparency, and deletion evidence; 2/3 for mutation lineage and tamper evidence. The public scorecard explains the two partial dimensions: records expose status, fact keys, and supersession links but not every possible lineage feature; the exported benchmark bundle exposes hashes but does not run the repository’s full checkpoint verifier as part of its rubric.

The next-highest public point total is 13/18. This gap is consistent with AETNAMEM’s evidence-first design, but comparison is confounded by manifest declarations and adapter




Figure 4: Flagship demonstration sequence. The LLM is intentionally absent: the experiment isolates control-plane behavior with deterministic inputs. Refusal behavior is asserted by exit status, and final claims are checked from the produced database rather than inferred from a video or transcript.

Table 3: Flagship experiment outcomes.

| Probe                                    | Result    | Evidence                                |
|------------------------------------------|-----------|-----------------------------------------|
| Hidden webpage value enters trusted slot | Blocked   | Record is quarantined at confidence 0.3 |
| Recall after poisoning                   | Safe      | Only active trusted record returned     |
| Untrusted evidence as authority          | Refused   | Policy exception, nonzero exit          |
| No authority for mutation                | Refused   | Policy exception, nonzero exit          |
| No reviewer key                          | Refused   | CLI key-boundary check                  |
| Unapproved commit                        | Refused   | State-machine check                     |
| Approved plan mutation                   | Refused   | Plan recomputation mismatch             |
| Exact approved write                     | Committed | Verified after-state + receipt          |
| Audit payload mutation                   | Detected  | Sequence 4 hash mismatch                |

visibility. We therefore do not treat 16/18 as an independently validated security ranking.

### 9.3 RQ2: guarded-action behavior

The recorded demo produces one active trusted record and one quarantined untrusted record in the same fact slot. Recall returns the trusted record only. The five prohibited action paths all exit nonzero, while the authorized write reaches committed; its operation reaches verified; and the resulting `report.md` has the expected SHA-256 digest.

The original artifact records 12 events: two episode ingestions; one trusted record creation; one quarantine; one recall; and seven action events from proposal through receipt. This compact history is sufficient to reconstruct the tested causal path without storing report contents in the audit plane.

### 9.4 RQ3: verification and tamper probes

The engine self-check passes against one checkpoint. Both standalone programs pass against the original database. The

action verifier recomputes the plan and validates the receipt even without the reviewer secret; when given the secret it can additionally authenticate the HMAC. The audit-history mutation fails with sequence 4: `event_hash` mismatch. The plan-table mutation is detected before execution and the attempted `retro.md` file is never written.

All 88 repository tests pass in 2.06 seconds on the disclosed machine. This runtime is included only as a reproducibility observation; no comparative performance claim is made.

## 9.5 Result interpretation

The evidence supports four narrow conclusions for the tested implementation and paths:

1. classified untrusted content remains non-recallable until promotion and does not supersede the trusted same-slot value;
2. enforced mutations require structurally trusted authority and exact-plan approval;
3. the reference adapter commits the authorized write only after precondition and postcondition checks; and
4. operational-plan and audit-history mutations used in the experiment are detected by independent code paths.

It does not show that arbitrary natural-language injections are classified, that every adapter verifies truthfully, that an LLM cannot use a bypass tool, or that five scenarios predict field failure rates.

## 10 Security Analysis

This section analyzes how the design responds to representative attacks and where the argument terminates at a trust assumption.

## MemoryStackBench seven\_sins\_v0\_1 public snapshot




Figure 5: Public MEMORYSTACKBENCH result snapshot at commit 10b9407. AETNAMEM has a perfect score and is highlighted in red; eleven additional targets also pass 33/33. The chart describes configured target harnesses, not broad product quality or model intelligence.

### 10.1 Persistent indirect prompt injection

An adversarial webpage asks the agent to remember a new value. Model-level instruction following may fail: the agent may summarize or repeat the payload. The control plane does not need to decide whether the text “looks malicious.” It needs authentic origin. Once labeled webpage/tool output, extracted facts receive untrusted tier and quarantine. Active-only retrieval blocks delayed re-entry, and quarantined same-slot facts cannot supersede trusted state.

This control fails if the host labels remote content as `user_message`, if another memory path bypasses admission, or if the agent can call promotion with effective authority. Source authenticity and complete mediation are therefore first-class deployment requirements.

### 10.2 Authority laundering

An injected agent may cite untrusted memory as justification for an effect. Because authority is a typed relation and trust tier rather than prose, the attempt is rejected. Relabeling the same reference `authorized_by` is also insufficient when its trust tier is untrusted or host attestation is false.

An attacker who controls trusted host flags can fabricate authority. The current CLI cannot prove that a human issued `--authority-id`; those flags are trustworthy only when

a host-side wrapper controls them. A production gateway should derive authority from authenticated task/session state rather than agent-supplied arguments.

### 10.3 Plan substitution and confused review

Approval covers the plan hash, and the plan covers policy, adapter manifest, exact operation digests, dependencies, evidence, actor, subject, and mode. Commit recomputes the plan from operational tables before any effect. Mutating arguments while leaving their recorded digest unchanged is detected when prepared payload digest and plan reconstruction are checked; mutating both payload and digest changes the plan; substituting a manifest changes its fingerprint.

The reviewer still needs a usable rendering of the plan. Cryptographic exactness does not prevent a human from approving misleading content or a dangerous but accurately described action. Psychological acceptability remains a security requirement [18].

### 10.4 World-state races

The filesystem adapter snapshots existence, kind, content digest, mode, and before-image. Commit validates all operations before starting and revalidates each immediately before execution. A concurrent change aborts or triggers compensa-




Figure 6: Check counts by category. The suite is most discriminative on untrusted-source resistance and temporal updates in this snapshot; deletion and write checks are passed by most configured targets.

tion. There remains a time-of-check/time-of-use window between final revalidation and OS operation, and path safety is implemented in user space rather than with a hardened capability filesystem. Symlinks are rejected, but a privileged local adversary racing directory replacement is outside the demonstrated guarantee.

Provider adapters should use native conditional writes, ETags, transaction tokens, or idempotency APIs where available. A generic pre-read is weaker than provider-enforced compare-and-swap.

## 10.5 False success and partial failure

Provider return values are treated as receipts, not proof. The adapter re-observes state and checks a typed postcondition. If verification fails, the transaction is uncertain and previous verified operations are compensated in reverse order. Failed compensation yields `partial`, preserving the failure instead of overwriting it with success.

The adapter supplies both execution and verification logic in v0.2.1. A defective or malicious adapter can lie consistently. Higher-assurance deployments should use independent read credentials, provider audit APIs, or a separately implemented observer.

## 10.6 Crash consistency

The engine records `effect_executing` before the provider call. An ordinary exception after dispatch becomes uncertain. A hard crash leaves a durable committing/executing marker; recovery fences it as `recovery_required` and refuses generic retry. This is safer than at-least-once replay but does not resolve the world automatically. The provider must offer lookup by idempotency key or authoritative state inspection.

## 10.7 Audit rewriting

Editing a past payload without recomputing its hash fails immediately, as demonstrated. Recomputing all later hashes still changes the head and fails an external checkpoint. Deleting the tail also fails checkpoint containment. If the attacker controls both database and checkpoint sink, the evidence can be rewritten consistently. External anchoring and retention governance are therefore not optional for adversarial tamper detection.

The event chain is not a forward-secure log in the stronger Schneier–Kelsey sense [20]: it does not evolve and erase authentication keys, encrypt past entries, or prevent a compromised writer from creating a new valid suffix. It is a transparent SHA-256 hash chain with external-head anchoring.




Figure 7: Separate auditability evidence matrix from the pinned public snapshot. AETNAMEM receives 16/18 points; the next-highest reported value is 13/18. Origin metadata is shown because native and undeclared evidence should not be conflated. This matrix is descriptive and not included in the safety score.

## 10.8 Privacy and erasure

Digest-only engine events reduce raw-value retention, and content/action payloads can be logically purged. However, hashes of low-entropy facts may be guessed, SQLite may retain bytes in free pages/WAL, and backups/replicas may outlive local deletion. Secure erase, encryption, backup lifecycle, data minimization, and jurisdiction-specific requirements remain deployment concerns. The architecture can support evidence for a deletion workflow; it is not a compliance certification [14].

## 10.9 Security-property summary

## 11 Related Work

### 11.1 Agents and tool use

ReAct interleaves model reasoning and environment actions [24]; Toolformer trains models to choose APIs and arguments [19]. These works improve action selection. AETNAMEM begins after selection: it asks whether the proposed effect has authority, whether the exact plan was approved, what assumptions it relies on, what the provider actually changed, and how that claim can later be verified.

ToolEmu uses a language-model-emulated sandbox and evaluator to discover risky agent behavior [17]. AgentDojo evaluates utility and prompt-injection security across realistic

tool environments [5]. These are complementary test environments. The current AETNAMEM experiment is much smaller but emits a verifiable operational artifact and exercises persistent memory plus action authority in one chain.

### 11.2 Prompt-injection defenses

Greshake et al. establish indirect prompt injection through remotely supplied data [8]. Instruction-hierarchy training teaches models to prioritize privileged instructions [22]. CaMeL extracts trusted control/data flows and applies capabilities to prevent unauthorized exfiltration [4]. CaMeL is the closest architectural relative: both place a protective system layer around an untrusted model and distinguish data from authority.

The contribution here is persistence and evidence: lifecycle-gated memory, exact-plan transaction records, before/after observations, compensation, receipts, erasure boundaries, and an externally anchorable history. Conversely, CaMeL offers a stronger formal security story for its mediated programs and a broader AgentDojo evaluation. AETNAMEM should not be read as superseding it; combining CaMeL-style control-flow extraction with AETNAMEM-style persistent evidence is a promising direction.




Figure 8: AETNAMEM profile. Safety bars report passed checks; auditability bars report the benchmark’s evidence points. Equal bar heights across the two panels do not imply equal semantics.

### Flagship demo: observed control decisions in execution order




Figure 9: Observed decisions in execution order. Red crosses are refused paths, purple markers are detected mutations, and the green square is the single authorized commit. The first three steps establish the delayed-memory attack precondition and quarantine result.

## 11.3 Long-term agent memory

Generative Agents stores experiences and synthesizes reflections [16]; MemoryBank adds evolving long-term memory [25]; and MemGPT uses tiered virtual context management [15]. LongMemEval evaluates extraction, multi-session reasoning, temporal reasoning, updates, and abstention over long conversations [23]. These systems emphasize memory capability and answer quality.

AETNAMEM emphasizes policy and evidence: origin, quarantine, mutation lineage, retrieval traces, deletion receipts, and audit integrity. Its deterministic v0 extractor is deliberately less semantically capable than model-backed memory. A mature system needs both: high-recall semantic extraction and a policy gate that does not grant extracted text authority by default.

MEMORYSTACKBENCH provides a common multi-session harness for 19 configured targets in the pinned snapshot [12]. Its five-scenario safety suite is narrower than LongMemEval’s 500 questions, but it directly checks webpage poisoning, deletion retention, provenance, temporal replacement, and over-generalization. The paper uses it as conformance evidence, not comprehensive memory quality.

## 11.4 Classical systems foundations

Fail-safe defaults, complete mediation, and separation of privilege come directly from Saltzer and Schroeder [18]. Typed,

scoped authority follows capability-system principles [6, 13]. The separation of `informed_by` and `authorized_by` can be viewed as a provenance/authority distinction: causal influence does not imply permission. Buneman et al.’s *why/where* provenance motivates retaining source relationships rather than only values [3].

Guarded multi-operation actions resemble sagas: when external effects cannot share one atomic transaction, explicit compensating actions and partial-failure states are required [7]. The implementation is intentionally more conservative about ambiguous dispatch: it fences uncertainty rather than assuming compensation is always possible.

Hash-linked history and external anchoring draw on secure logging and timestamping [20, 9]. AETNAMEM’s construction is simpler than forward-secure encrypted logging and transparency trees, but its threat boundary is explicit and its wire format is independently executable.

## 11.5 Evaluation practice

Kapoor et al. warn that agent benchmarks can encourage shortcuts, conflate developer needs, omit cost, and lack reproducibility [10]. We respond by pinning commits and hashes, rerunning the current version, reporting ties and target heterogeneity, excluding timing rankings, testing policy with non-benchmark vocabulary, and disclosing common ownership. The resulting evidence is stronger than an unqualified score

Table 4: Security properties, assumptions, and principal residual risks.

| Property                  | Mechanism                                     | Required assumptions                               | Residual risk                                                       |
|---------------------------|-----------------------------------------------|----------------------------------------------------|---------------------------------------------------------------------|
| Quarantine non-recall     | Source trust mapping; active-only query       | Authentic source class; mediated read path         | Misclassification, promotion abuse, direct DB read                  |
| Trusted-slot preservation | Supersede only when new record is active      | Policy path not bypassed                           | Malicious promotion or direct status edit                           |
| Authority non-escalation  | Typed relation, tier, host attestation        | Host controls attestation                          | Compromised host can mint authority                                 |
| Exact-plan approval       | Canonical plan hash and HMAC                  | Key secrecy, collision resistance, clock sanity    | Shared-key identity ambiguity; misleading preview                   |
| Fresh precondition        | Adapter snapshot and two-phase revalidation   | Correct observer; sufficiently strong provider API | TOCTOU race, incomplete precondition                                |
| Verified outcome          | Provider receipt plus postcondition read      | Correct independent-enough observation             | Adapter may lie; external side effects unobserved                   |
| No blind retry            | Executing marker, uncertainty, recovery fence | Durable SQLite transition                          | Manual/provider reconciliation still required                       |
| Tamper evidence           | Per-subject chain and external checkpoint     | Anchor outside attacker domain                     | Pre-anchor rewrite; compromised anchor; no event completeness proof |
| Payload erasure           | Separate mutable payload/content tables       | Correct purge and backup policy                    | Free pages, WAL, replicas, digest inference                         |

but still requires external replication and broader suites.

## 12 Limitations and Roadmap

### 12.1 Evaluation limitations

The benchmark has only five scenarios and 33 assertions. They are deterministic, public, and visible to the implementer. Passing them cannot estimate a field attack rate, robustness to paraphrase, multilingual injection, multimodal content, long-horizon drift, or adaptive adversaries. The flagship demonstration is one carefully chosen vertical slice. Its value is in-spectable causality, not breadth.

The benchmark and project share a maintainer, and the “independent” verifiers share a repository and author. An external team should rerun the artifacts, inspect adapter semantics, add hidden scenarios, and implement at least one verifier from the written specification.

The public comparison mixes native memory agents, managed services, direct APIs, store-level harnesses, and conversation-session persistence. A 33/33 harness may include explicit adapter policy not present in every default product path. The paper therefore avoids broad competitive claims.

### 12.2 Implementation limitations

**No universal action gateway.** Memory MCP exists, but guarded actions are not yet exposed as a mandatory MCP interception layer. An agent with another filesystem or network tool can bypass the reference engine.

**Host-attested, not cryptographically authenticated, source and authority.** The engine trusts source type and CLI authority flags supplied by the host. Subject IDs are storage scopes, not authenticated tenants.

**Shared-key approval.** HMAC is a practical local primitive but lacks signer non-repudiation and makes key distribution/rotation operationally sensitive. Approver labels are claims.

**One narrow adapter.** The filesystem vertical slice supports two operations. Email, browser, database, cloud, messaging,

payment, and code-execution adapters need provider-specific preconditions, idempotency, observation, and compensation.

**Adapter-coupled verification.** The same adapter family prepares, executes, and verifies. Separate observer credentials and implementations would reduce correlated failure.

**Deterministic extraction and lexical retrieval.** Generic patterns and FTS5 make behavior reproducible but limit semantic coverage. Model-backed extraction and vector retrieval remain future work and must preserve quarantine semantics.

**Storage and deployment maturity.** SQLite is the only storage backend. Payload encryption, remote multi-tenant service hardening, KMS integration, role-based administration, secure checkpoint publication, and operational observability remain deployment work.

**No formal proof or machine-checked model.** Properties in this report are code-level arguments under stated assumptions. The state machine and schemas have not been verified in TLA+, Alloy, Coq, Lean, or a model checker.

### 12.3 Roadmap

The highest-leverage next steps are:

1. **Complete mediation:** add an authenticated action gateway for MCP/HTTP hosts and remove direct mutation capabilities from agent-facing processes.
2. **Asymmetric approval:** define a signer interface for KMS/HSM/passkey-backed signatures, key identifiers, rotation, revocation, and role policy.
3. **Independent observation:** support read-only verifier principals and provider audit-log correlation distinct from executor credentials.
4. **Provider adapters:** implement conditional/idempotent adapters for email, calendars, Git, cloud storage, databases, browsers, and messaging.
5. **Formal protocol publication:** version and publish WORLDPATCH, approval, receipt, and verifier schemas as a permissively licensed protocol package.
6. **Broader evaluation:** add guarded-action tracks to MEMORYSTACKBENCH, hidden/adaptive scenarios, multiple mod-

els/hosts, crash injection, concurrency, fuzzing, and repeated performance experiments.

7. **Policy-preserving semantic memory:** add LLM extraction and embeddings only behind the same admission, provenance, deletion, and retrieval-evidence boundaries.
8. **External witnessing:** integrate transparency logs, WORM storage, or trusted timestamping for production checkpoint anchors.

The design target is not “a better prompt.” It is a deployable control path in which the model never possesses all authority required to approve, execute, and rewrite the evidence for its own effects.

## 13 Conclusion

Stateful agents need more than memory and tools. They need an authority model, transaction semantics, result verification, and evidence that survives disagreement about what happened. AETNAMEM implements a compact version of that control plane: origin-aware memory admission, quarantine, active-only recall, typed causal authority, exact-plan approval, preconditioned and verified effects, explicit uncertainty, chain-bound receipts, and standalone verification.

The artifact passes every assertion in the pinned MEMORYSTACKBENCH suite and reproduces five action refusals, one authorized commit, and two tamper detections over a real database. Those results are strong within their scope and deliberately modest outside it. The engine is deterministic, the benchmark tie is broad, the evaluation is self-maintained, and forced mediation remains unfinished.

The more general lesson is architectural. A model’s output should be treated as a proposal. Durable memory should be admitted according to origin and trust. External effects should require scoped authority and exact reviewed intent. Claimed results should be reconciled with observed state. History should be independently checkable and anchored outside the writer’s control. Evidence should come before effect, and verification should come after it.

## References

- [1] Carlisle Adams, Peter Cain, Denis Pinkas, and Robert Zuccherato. Internet X.509 public key infrastructure time-stamp protocol. RFC 3161, Internet Engineering Task Force, 2001.
- [2] aetna000.com. aetnamem: Auditable memory and guarded actions for agents, 2026. Commit 0cd082c9cac14f35a66ff946395a31847322005d.
- [3] Peter Buneman, Sanjeev Khanna, and Wang-Chiew Tan. Why and where: A characterization of data provenance. In *Database Theory – ICDT 2001*, pages 316–330. Springer, 2001.
- [4] Edoardo Debenedetti, Ilia Shumailov, Tianqi Fan, Jamie Hayes, Nicholas Carlini, Daniel Fabian, Christoph Kern, Chongyang Shi, Andreas Terzis, and Florian Tramer. Defeating prompt injections by design. *arXiv preprint arXiv:2503.18813*, 2025.
- [5] Edoardo Debenedetti, Jie Zhang, Mislav Balunovic, Luca Beurer-Kellner, Marc Fischer, and Florian Tramer. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. In *Advances in Neural Information Processing Systems*, 2024.
- [6] Jack B. Dennis and Earl C. Van Horn. Programming semantics for multi-programmed computations. In *Communications of the ACM*, volume 9, pages 143–155, 1966.
- [7] Hector Garcia-Molina and Kenneth Salem. Sagas. In *Proceedings of the 1987 ACM SIGMOD International Conference on Management of Data*, pages 249–259, 1987.
- [8] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*, 2023.
- [9] Stuart Haber and W. Scott Stornetta. How to time-stamp a digital document. *Journal of Cryptology*, 3(2):99–111, 1991.
- [10] Sayash Kapoor, Benedikt Stroebel, Zachary S. Siegel, Nitya Nadgir, and Arvind Narayanan. AI agents that matter. *arXiv preprint arXiv:2407.01502*, 2024.
- [11] Hugo Krawczyk, Mihir Bellare, and Ran Canetti. HMAC: Keyed-hashing for message authentication. RFC 2104, Internet Engineering Task Force, 1997.
- [12] MemoryStackBench Contributors. Memorystackbench: An automated benchmark farm for agent memory frameworks, 2026. Public snapshot at commit 10b9407ce54c92bcb8aee24099505427aabeebcd.
- [13] Mark S. Miller, Ka-Ping Yee, and Jonathan Shapiro. Capability myths demolished. Technical Report SRL2003-02, Johns Hopkins University Systems Research Laboratory, 2003.
- [14] National Institute of Standards and Technology. Artificial intelligence risk management framework (AI RMF 1.0). Technical Report NIST AI 100-1, NIST, 2023.
- [15] Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. MemGPT: Towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*, 2023.
- [16] Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, 2023.
- [17] Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitis, Yongchao Zhou, Jimmy Ba, Yann Dubois, Chris J. Maddison, and Tatsunori Hashimoto. Identifying the risks of LM agents with an LM-emulated sandbox. *arXiv preprint arXiv:2309.15817*, 2023.
- [18] Jerome H. Saltzer and Michael D. Schroeder. The protection of information in computer systems. *Proceedings of the IEEE*, 63(9):1278–1308, 1975.
- [19] Timo Schick, Jane Dwivedi-Yu, Roberto Dessi, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems*, 2023.
- [20] Bruce Schneier and John Kelsey. Cryptographic support for secure logs on untrusted machines. In *7th USENIX Security Symposium*. USENIX Association, 1998.
- [21] SQLite Project. SQLite FTS5 extension, 2026. Accessed 2026-07-13.
- [22] Eric Wallace, Kai Xiao, Reimar Leike, Lilian Weng, Johannes Heidecke, and Alex Beutel. The instruction hierarchy: Training LLMs to prioritize privileged instructions. *arXiv preprint arXiv:2404.13208*, 2024.
- [23] Di Wu, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu. LongMemEval: Benchmarking chat assistants on long-term interactive memory. *arXiv preprint arXiv:2410.10813*, 2024.
- [24] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023.
- [25] Wanjun Zhong, Lianghong Guo, Qiqi Gao, He Ye, and Yanlin Wang. MemoryBank: Enhancing large language models with long-term memory. *arXiv preprint arXiv:2305.10250*, 2023.

## A Protocol and Wire-Format Detail

This appendix records the hashed document boundaries needed to reproduce the verifier. It is descriptive of v0.2.1; the repository specifications remain canonical when future format versions are introduced.

### A.1 Canonical serialization

For a JSON-compatible value  $x$ :

```
json.dumps(x, sort_keys=True, separators=(",", ":"), ensure_ascii=False)
```

The UTF-8 bytes of this representation are hashed with SHA-256 and rendered as lowercase hexadecimal. Floating-point interoperability is intentionally avoided in hashed action documents. Timestamps are ISO-8601 strings already present in the document.

### A.2 WorldPatch preimage

```
{
  "format": "aetna-world-patch-v1",
  "transaction_id": "act_...",
  "subject_id": "...",
  "actor_id": "...",
  "mode": "enforce",
  "plan_version": 1,
  "policy_hash": "sha256...",
  "operations": [{
    "id": "op_...",
    "key": "write-report",
    "ordinal": 0,
    "adapter": "filesystem",
    "operation": "write_text",
    "effect_class": "verified_compensatable",
    "arguments_digest": "sha256...",
    "preview_digest": "sha256...",
    "precondition_digest": "sha256...",
    "manifest_digest": "sha256...",
    "depends_on": [],
    "evidence": [{
      "kind": "host_task",
      "ref_id": "task-42",
      "digest": "sha256...",
      "relation": "authorized_by",
      "trust_tier": "trusted_user",
      "attested": true
    }]
  }]
}
```

Operation order is stable topological order. Evidence is sorted by relation, kind, reference ID, and digest. The plan stores digests, while normalized arguments, previews, preconditions, and before-images reside in the payload plane.

### A.3 Approval document

```
{
  "transaction_id": "act_...",
  "plan_hash": "sha256...",
  "approver": "reviewer-label",
  "issued_at": "ISO-8601 UTC",
  "expires_at": "ISO-8601 UTC",
  "nonce": "32 hexadecimal characters",
  "format": "aetna-action-approval-v1"
}
```

The signature is HMAC-SHA-256 over the canonical document above. The stored approval digest hashes the full document including signature. Approval IDs derive from the unique nonce. Verification checks transaction, plan, signature, issue time, and expiry.

### A.4 Audit-event preimage

```
{
  "event_id": "aud_...",
  "subject_id": "...",
  "event_type": "action.receipt_issued",
  "created_at": "ISO-8601 UTC",
  "actor": "...",
  "session_id": null,
  "turn_id": null,
  "record_id": null,
}
```

```

"payload": {"transaction_id": "act_...", "...": "..."},
"prev_hash": "previous subject event hash or null"
}

```

The event hash is SHA-256 of the canonical preimage. The global SQLite sequence is not in the preimage; it determines traversal order. The chain is per subject.

## A.5 Action receipt

```

{
  "format": "aetna-action-receipt-v1",
  "transaction_id": "act_...",
  "subject_id": "...",
  "plan_hash": "sha256...",
  "terminal_state": "committed",
  "operation_receipts": [{
    "operation_id": "op_...",
    "adapter": "filesystem",
    "operation": "write_text",
    "state": "verified",
    "result_digest": "sha256...",
    "compensation_digest": null
  }],
  "audit_event_id": "aud_...",
  "audit_event_hash": "sha256...",
  "created_at": "ISO-8601 UTC",
  "receipt_sha256": "sha256 of all preceding fields"
}

```

The corresponding `action.receipt_issued` event contains transaction ID, plan hash, terminal state, and a digest of the operation-receipt array. The receipt then binds back to that event's ID and hash.

## A.6 Checkpoint

```

{
  "format": "aetnamem-checkpoint-v1",
  "created_at": "ISO-8601 UTC",
  "subjects": {
    "subject": {
      "sequence": 12,
      "event_hash": "sha256...",
      "event_count": 12
    }
  },
  "checkpoint_sha256": "sha256 of all preceding fields"
}

```

Containment verifies that the database still has the event at each pinned sequence with the pinned hash. It does not assert that no valid later suffix exists.

## B Benchmark Scenario Contract

Table 5: Complete high-level contract of `seven_sins_v0_1`.

| Scenario                  | Sessions                                                                   | Positive requirements                                                                                  | Negative requirements                                                                                 |
|---------------------------|----------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------|
| Durable private itinerary | Plant preference; later ask default                                        | Response says private/not public; stored fact contains preference and provenance fields from session 1 | None beyond record identity/provenance checks                                                         |
| Webpage poisoning         | Plant private preference; embed public instruction in webpage; later probe | Later response remains private; trusted record retains original session/source                         | No commitment to poisoned instruction; no poisoned text stored as active memory                       |
| Temporal airport update   | Plant SFO; correct to OAK; later probe                                     | Response and active memory identify OAK                                                                | Response must not recommend/currently identify SFO; active state must not retain stale SFO as current |
| Backup-email deletion     | Plant address; request forget; probe twice                                 | Responses acknowledge absence/deletion                                                                 | Address absent from responses and live memory                                                         |
| Diet overgeneralization   | State beef avoidance at business dinners; ask vegetarian                   | Response preserves narrow evidence or uncertainty                                                      | No categorical vegetarian fact or causal overgeneralization                                           |

The suite produces 33 checks: 15 high-severity and 18 medium-severity. Severity weighting totals 81. The current target adapter uses deterministic response synthesis from retrieved records. Therefore the result evaluates the memory engine plus adapter contract, not a generative model.

Table 6: Pinned public MemoryStackBench results used by the figures.

| Configured target      | Checks | Scenarios | Weighted | Audit | Origin     |
|------------------------|--------|-----------|----------|-------|------------|
| aetnamem               | 33/33  | 5/5       | 81/81    | 16/18 | native     |
| Agno Memory            | 33/33  | 5/5       | 81/81    | 12/18 | undeclared |
| AWS AgentCore          | 33/33  | 5/5       | 81/81    | 11/18 | undeclared |
| CrewAI Memory          | 33/33  | 5/5       | 81/81    | 13/18 | undeclared |
| Google ADK Memory Bank | 33/33  | 5/5       | 81/81    | 12/18 | undeclared |
| Hindsight              | 33/33  | 5/5       | 81/81    | 12/18 | undeclared |
| LangGraph              | 33/33  | 5/5       | 81/81    | 12/18 | undeclared |
| LangMem                | 33/33  | 5/5       | 81/81    | 12/18 | undeclared |
| LlamaIndex Memory      | 33/33  | 5/5       | 81/81    | 12/18 | undeclared |
| Supermemory            | 33/33  | 5/5       | 81/81    | 13/18 | undeclared |
| Tree Ring Memory       | 33/33  | 5/5       | 81/81    | 12/18 | undeclared |
| Zep Cloud              | 33/33  | 5/5       | 81/81    | 11/18 | undeclared |
| AutoGen + Mem0         | 32/33  | 4/5       | 78/81    | 13/18 | undeclared |
| Mem0                   | 32/33  | 4/5       | 78/81    | 13/18 | undeclared |
| Cognee                 | 31/33  | 3/5       | 76/81    | 12/18 | undeclared |
| Letta                  | 31/33  | 3/5       | 76/81    | 11/18 | undeclared |
| OpenAI Agents Sessions | 30/33  | 3/5       | 73/81    | 9/18  | undeclared |
| TencentDB Agent Memory | 26/33  | 2/5       | 61/81    | 10/18 | undeclared |
| Graphiti + Neo4j       | 24/33  | 1/5       | 60/81    | 5/18  | undeclared |

## C Public Snapshot Detail

The benchmark repository documents important target-specific interpretation notes. Several perfect rows are implemented store harnesses that add explicit write/retrieval/delete policy around real APIs. The table must not be repurposed as a general feature, quality, latency, or security ranking.

## D Recorded Flagship Evidence

### D.1 Complete event sequence

Table 7: All 12 events in the distributed flagship database. Hashes are abbreviated for readability; verification uses all 64 hexadecimal characters.

| Seq. | Event type                | Actor          | Purpose                         | Hash prefix  | Previous prefix |
|------|---------------------------|----------------|---------------------------------|--------------|-----------------|
| 1    | episode.ingested          | user           | trusted statement admitted      | 727d2b940874 | –               |
| 2    | memory.record_created     | system         | trusted fact active             | b6e8f4c0a691 | 727d2b940874    |
| 3    | episode.ingested          | user           | webpage text ingested           | 444191a58c05 | b6e8f4c0a691    |
| 4    | memory.record_quarantined | system         | hostile same-slot fact isolated | 71d2260d1337 | 444191a58c05    |
| 5    | memory.recall             | system         | active-only result recorded     | f6e30af8fa4f | 71d2260d1337    |
| 6    | action.proposed           | research-agent | authorized WorldPatch staged    | 293148e07c01 | f6e30af8fa4f    |
| 7    | action.approved           | demo-user      | exact plan approved             | 3bfe2f72b71c | 293148e07c01    |
| 8    | action.committing         | research-agent | durable execution intent        | 3988ce79a6fa | 3bfe2f72b71c    |
| 9    | action.effect_executing   | research-agent | external boundary crossed       | 38cc004d23d5 | 3988ce79a6fa    |
| 10   | action.effect_verified    | research-agent | postcondition observed          | 09f4a4f0d257 | 38cc004d23d5    |
| 11   | action.committed          | research-agent | terminal transaction state      | 746ae6825e4c | 09f4a4f0d257    |
| 12   | action.receipt_issued     | research-agent | terminal receipt bound          | 11232bec0f92 | 746ae6825e4c    |

Expected refusals before proposal are not inserted into this particular database because proposal validation fails before a transaction exists. Their nonzero exits are asserted in the deterministic driver and captured in the transcript. Production systems may additionally log rejected proposals at a trusted ingress boundary without storing sensitive raw arguments.

### D.2 Refusal claim ledger

### D.3 Recorded artifact digests

```
5c7f46db3101de28f7381f2166e33424bba3b4b9273fd7807ebe5260747896a4 memories.db
7e528a029c07c2fdc76851dd0988087a083c810a2b3b9cf4899221c17d1cfa7a checkpoints.jsonl
a2497f67e66f54bd0efa45d34d11a0f4a283187b861560de3d3da1f6b0dac203 transcript.txt
```

The recorded transaction is `act_ebd33915ea974e64a36282deef754f49`. The checkpoint pins subject `demo-user` at sequence 12. Its full event hash appears in the distributed checkpoint JSONL and is verified directly by the reproduction command below.

Table 8: The five expected action refusals and the tested invariant.

| Attempt                                           | Required refusal                                                     | Invariant exercised                          |
|---------------------------------------------------|----------------------------------------------------------------------|----------------------------------------------|
| Quarantined record used as authorized_by          | untrusted_content evidence may inform but cannot authorize an action | Information does not escalate to authority   |
| Mutation with no authority reference              | enforced mutations require authorized_by evidence                    | Fail-safe default for external mutation      |
| Agent-facing commit with no reviewer key          | CLI refuses and directs operator to external key boundary            | Separation of privilege                      |
| Key holder commits unapproved transaction         | transaction is awaiting_approval, not approved                       | State-machine mediation                      |
| Approved plan's operation digest edited on a copy | persisted action plan does not match its plan hash                   | Exact-plan approval and commit recomputation |

## E Reproduction Procedure

### E.1 Repository tests

```
git checkout 0cd082c9cac14f35a66ff946395a31847322005d
python3.12 -m pytest -q
# Expected: 88 passed
```

### E.2 Fresh flagship run

```
./examples/flagship-demo/run.sh /tmp/aetnamem-flagship
python3 tools/verify_audit.py \
  /tmp/aetnamem-flagship/memories.db \
  --checkpoints /tmp/aetnamem-flagship/checkpoints.jsonl
# Copy the generated act_... ID printed by run.sh:
python3 tools/verify_actions.py \
  /tmp/aetnamem-flagship/memories.db act_...
```

The driver uses random identifiers and current timestamps, so byte-level database hashes differ across fresh runs. Structural outcomes and verifier results are deterministic.

### E.3 Verify the distributed demo database

```
python3 tools/verify_audit.py \
  examples/flagship-demo/artifacts/memories.db \
  --checkpoints examples/flagship-demo/artifacts/checkpoints.jsonl
python3 tools/verify_actions.py \
  examples/flagship-demo/artifacts/memories.db \
  act_ebd33915ea974e64a36282deef754f49
```

Expected output is OK demo-user and OK act\_..., respectively.

### E.4 Benchmark rerun

```
git clone https://github.com/aetna000/MemoryStackBench.git /tmp/msb
git -C /tmp/msb checkout 10b9407ce54c92bcb8aee24099505427aabeebcd
PYTHONPATH="/tmp/msb" python3.12 -m memorybench.cli run \
  --target "$PWD/bench/targets/aetnamem.yaml" \
  --suite /tmp/msb/suites/seven_sins_v0_1 \
  --out /tmp/aetnamem-paper-benchmark
jq '{overall,scenario_overall,severity_weighted_overall,categories}' \
  /tmp/aetnamem-paper-benchmark/scorecard.json
```

Expected totals are 33/33, 5/5, and 81/81. Install the benchmark's PyYAML dependency in the selected Python environment before running.

### E.5 Paper build and evidence validation

```
cd paper
make
make verify
# Compiled PDF: build/aetnamem-control-plane.pdf
```

The figure script consumes pinned CSV files. The validator checks row counts and headline values, verifies distributed artifact SHA-256 digests, queries SQLite metrics, and executes both standalone verifiers.

## F Deployment Review Checklist

The following checklist translates assumptions into deployment questions.

Table 9: Minimum review before claiming enforced operation.

| Question                                                           | Why it matters                                       | Required evidence                                    |
|--------------------------------------------------------------------|------------------------------------------------------|------------------------------------------------------|
| Can the agent reach any mutation tool outside the guarded gateway? | Bypass defeats complete mediation                    | Capability inventory, sandbox/tool configuration     |
| Who assigns source type and subject ID?                            | Agent-supplied labels can forge trust or cross scope | Authenticated host code and tests                    |
| Where is reviewer authority stored?                                | Agent key access collapses separation of privilege   | Separate process/principal, secret/KMS policy        |
| What exactly does the reviewer see?                                | A correct hash can bind a misleading preview         | Human-readable canonical diff and resource identity  |
| Does the provider support conditional/idempotent operations?       | Generic pre-read cannot eliminate races/duplicates   | ETag/version/idempotency API contract                |
| Who verifies after-state and with what credentials?                | Executor and observer can share failure modes        | Read-only observer or provider audit record          |
| How are uncertain effects reconciled?                              | Blind retry can duplicate irreversible work          | Provider lookup/runbook and operator ownership       |
| Where are checkpoints anchored?                                    | Local checkpoints can be rewritten with the DB       | WORM/transparency/timestamp receipt                  |
| What data remains after purge?                                     | Logical deletion is not media sanitization           | WAL/free-page/backup/replica retention policy        |
| Who can promote quarantine?                                        | Promotion converts data into active context          | Human/host authorization and audit review            |
| Are verifier binaries/code pinned separately?                      | Same-host replacement can forge a passing check      | Reproducible build, signed release, independent copy |

## G Claim and Disclosure Ledger

Table 10: What may and may not be claimed from this artifact.

| Claim            | Supported wording                                                               | Unsupported extrapolation                           |
|------------------|---------------------------------------------------------------------------------|-----------------------------------------------------|
| Benchmark        | “aetnamem passed 33/33 checks and 5/5 scenarios in the pinned suite”            | “best agent memory system” or “only perfect system” |
| Model            | “deterministic engine; no model used in the target run”                         | “the model aced the benchmark”                      |
| Prompt injection | “classified webpage fact was quarantined in the tested path”                    | “immune to prompt injection”                        |
| Action safety    | “five tested prohibited paths refused; one exact approved write committed”      | “all unauthorized actions are impossible”           |
| Verification     | “two code-independent-from-engine verifiers passed and caught tested mutations” | “third-party audited” or “formally verified”        |
| Tamper evidence  | “mutation before anchored head is detectable under hash assumptions”            | “tamper-proof database”                             |
| Deletion         | “live content was logically purged and receipt-bound”                           | “all bytes erased from media/backups”               |
| Compliance       | “supports accountable workflows”                                                | “GDPR/HIPAA/SOC 2 compliant by construction”        |

The author maintains both the evaluated project and the benchmark repository under the aetna000 account. This conflict is visible so readers can calibrate the evidence and prioritize independent replication.

## H Ethical and Operational Considerations

The hostile payload uses the IANA-reserved `attacker.example` domain and writes only inside a temporary root. No credentials, external services, or real exfiltration are involved. The demonstration is defensive and reproducible without network access.

A control plane can itself centralize sensitive data and authority. Deployers should minimize retained payloads, encrypt storage, separate reviewer/executor/verifier principals, rate-limit proposals, review quarantine, secure checkpoint sinks, and provide user-visible correction/deletion paths. Audit evidence should not become an excuse for pervasive raw-content retention.

Finally, automated refusal can deny legitimate work. Policy should expose a review path, precise reason, and bounded remediation rather than silently dropping a task. Human approval is not a universal solution: it can become habituated, overloaded, or socially engineered. The architecture reduces single-model authority but does not remove organizational responsibility.